

Biochemical Oxygen Demand

29-Oct-24

Name: Faizan

Registration # FA21-BSE-011

Dataset Assigned

BOD (Biochemical Oxygen Demand).

GitHub Link of Implementation

<https://github.com/Faizan-Kh/Data-Science>

Minkowski Distance

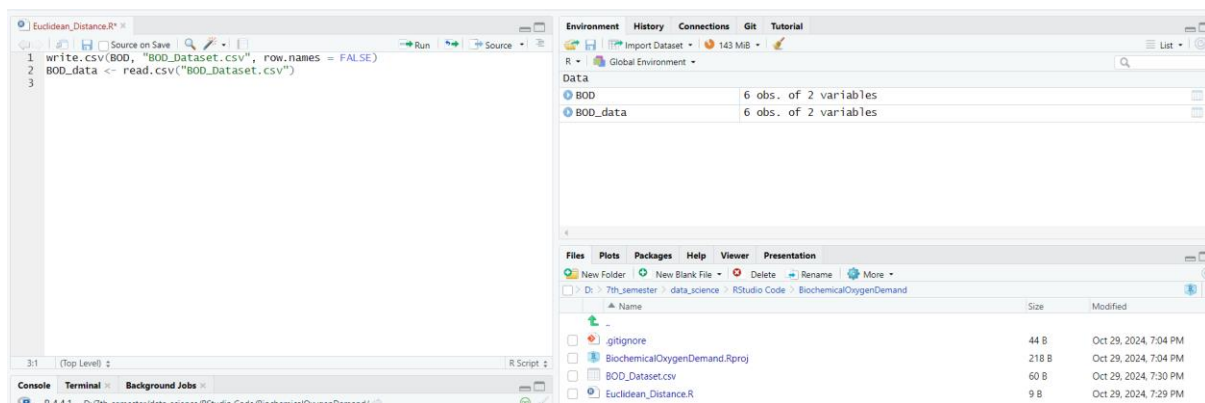
Minkowski distance is the generalization of Euclidean distance. Euclidean Distance Length of the line segment between two points in a Euclidean space is known as the Euclidean Distance [1].

RStudio implementation

1) Loading the dataset to memory:

```
> data("BOD")
> show(BOD)
  Time demand
1    1    8.3
2    2   10.3
3    3   19.0
4    4   16.0
5    5   15.6
6    7   19.8
```

2) Exporting the dataset to a CSV file and importing it



3) Calculating the Euclidean, city block and supremum distance by putting $r=1, 2$, and ∞ on the imported dataset

Load the BOD dataset and save it to a CSV

```
write.csv(BOD, "BOD_Dataset.csv", row.names = FALSE)
```

```

# Read the dataset back into R
BOD_data <- read.csv("BOD_Dataset.csv")

# Define a generalized Minkowski distance function
minkowski_dist <- function(x, y, r) {
  if (r == Inf) {
    return(max(abs(x - y))) #  $L^\infty$  norm (Supremum distance)
  } else {
    return((sum(abs(x - y)^r))^(1/r)) #  $L_r$  norm (Minkowski distance)
  }
}

# Get the number of demands
n <- length(BOD_data$demand)

# Initialize matrices to hold the distances
manhattan_distances <- matrix(NA, n, n)
euclidean_distances <- matrix(NA, n, n)
supremum_distances <- matrix(NA, n, n)

# Set values for r
r_values <- c(1, 2, Inf) # 1 = Manhattan, 2 = Euclidean, Inf = Supremum

# Calculate the distances for each pair
for (i in 1:n) {
  for (j in 1:n) {
    manhattan_distances[i, j] <-
minkowski_dist(BOD_data$demand[i], BOD_data$demand[j],
r_values[1])

```

```

        euclidean_distances[i, j] <-
minkowski_dist(BOD_data$demand[i], BOD_data$demand[j],
r_values[2])

        supremum_distances[i, j] <-
minkowski_dist(BOD_data$demand[i], BOD_data$demand[j],
r_values[3])
    }
}

```

```

# Print the distance matrices
cat("Manhattan Distance Matrix (r=1):\n")
print(manhattan_distances)
cat("\nEuclidean Distance Matrix (r=2):\n")
print(euclidean_distances)
cat("\nSupremum Distance Matrix (r=∞):\n")
print(supremum_distances)
print(distances)

```

Output

```
> # Print the distance matrices
> cat("Manhattan Distance Matrix (r=1):\n")
Manhattan Distance Matrix (r=1):
> print(manhattan_distances)
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]  0.0  2.0 10.7  7.7  7.3 11.5
[2,]  2.0  0.0  8.7  5.7  5.3  9.5
[3,] 10.7  8.7  0.0  3.0  3.4  0.8
[4,]  7.7  5.7  3.0  0.0  0.4  3.8
[5,]  7.3  5.3  3.4  0.4  0.0  4.2
[6,] 11.5  9.5  0.8  3.8  4.2  0.0
> cat("\nEuclidean Distance Matrix (r=2):\n")

Euclidean Distance Matrix (r=2):
> print(euclidean_distances)
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]  0.0  2.0 10.7  7.7  7.3 11.5
[2,]  2.0  0.0  8.7  5.7  5.3  9.5
[3,] 10.7  8.7  0.0  3.0  3.4  0.8
[4,]  7.7  5.7  3.0  0.0  0.4  3.8
[5,]  7.3  5.3  3.4  0.4  0.0  4.2
[6,] 11.5  9.5  0.8  3.8  4.2  0.0
> cat("\nSupremum Distance Matrix (r=∞):\n")

Supremum Distance Matrix (r=∞):
> print(supremum_distances)
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]  0.0  2.0 10.7  7.7  7.3 11.5
[2,]  2.0  0.0  8.7  5.7  5.3  9.5
[3,] 10.7  8.7  0.0  3.0  3.4  0.8
[4,]  7.7  5.7  3.0  0.0  0.4  3.8
[5,]  7.3  5.3  3.4  0.4  0.0  4.2
[6,] 11.5  9.5  0.8  3.8  4.2  0.0
> |
```

4)

SMC (Simple Matching Coefficient)

Also known as the statistic used for comparing the similarities between sample sets [2].

Implementation

- 1) We initialize a threshold to convert the data into binary

```
threshold <- mean(BOD_data$demand)
```

```
#Converting the data to binary
```

```
binary_demand <- ifelse(BOD_data$demand > threshold, 1, 0)
```

- 2) Treating both 1s and 0s

```
n <- length(binary_demand)
```

```
P_11 <- sum(binary_demand %>% t(binary_demand)) # Both 1s
```

```
P_00 <- sum((1 - binary_demand) %>% t(1 - binary_demand)) #Both 0s
```

- 3) Calculating the total number of similarities and SMC

```
#Total number of comparisons
```

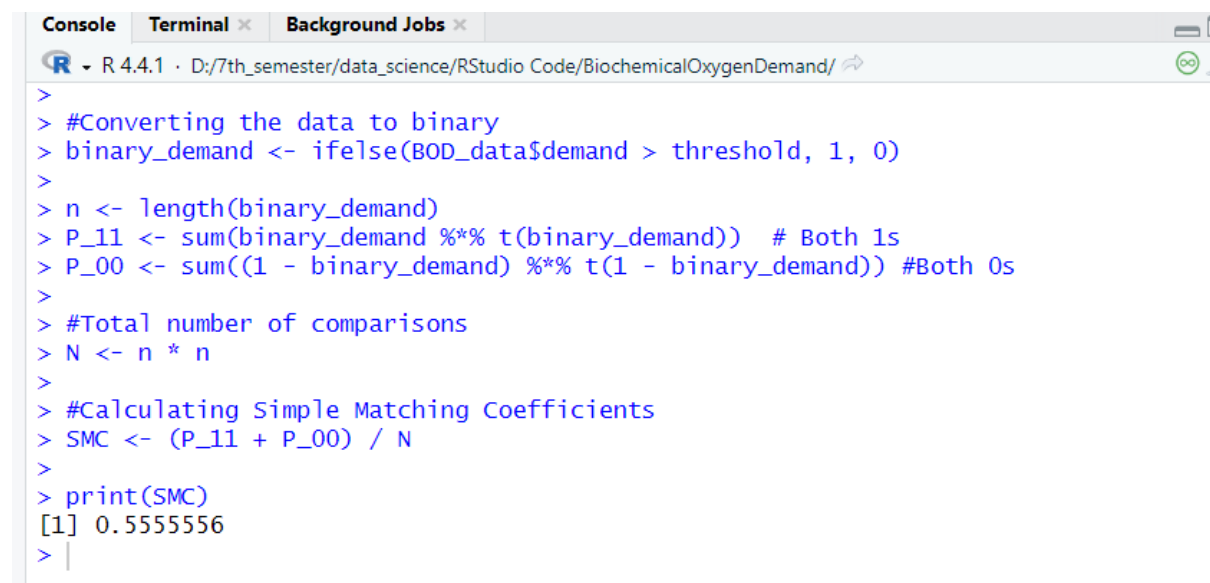
```
N <- n * n
```

```
#Calculating Simple Matching Coefficients
```

```
SMC <- (P_11 + P_00) / N
```

```
print(SMC)
```

Output



```
R - R 4.4.1 · D:/7th_semester/data_science/RStudio Code/BiochemicalOxygenDemand/
> 
> #Converting the data to binary
> binary_demand <- ifelse(BOD_data$demand > threshold, 1, 0)
> 
> n <- length(binary_demand)
> P_11 <- sum(binary_demand %>% t(binary_demand)) # Both 1s
> P_00 <- sum((1 - binary_demand) %>% t(1 - binary_demand)) #Both 0s
> 
> #Total number of comparisons
> N <- n * n
> 
> #Calculating Simple Matching Coefficients
> SMC <- (P_11 + P_00) / N
> 
> print(SMC)
[1] 0.5555556
> |
```

Jaccard Coefficient

Also known as Jaccard index, is a statistic used to find out the similarity and diversity of sample sets [3].

Implementation

```
threshold <- 15 # Example threshold

binary_demand <- ifelse(BOD$demand > threshold, 1, 0)


n <- length(binary_demand)


# Initialize counts for M11, M01, and M10

M_11 <- 0
M_01 <- 0
M_10 <- 0

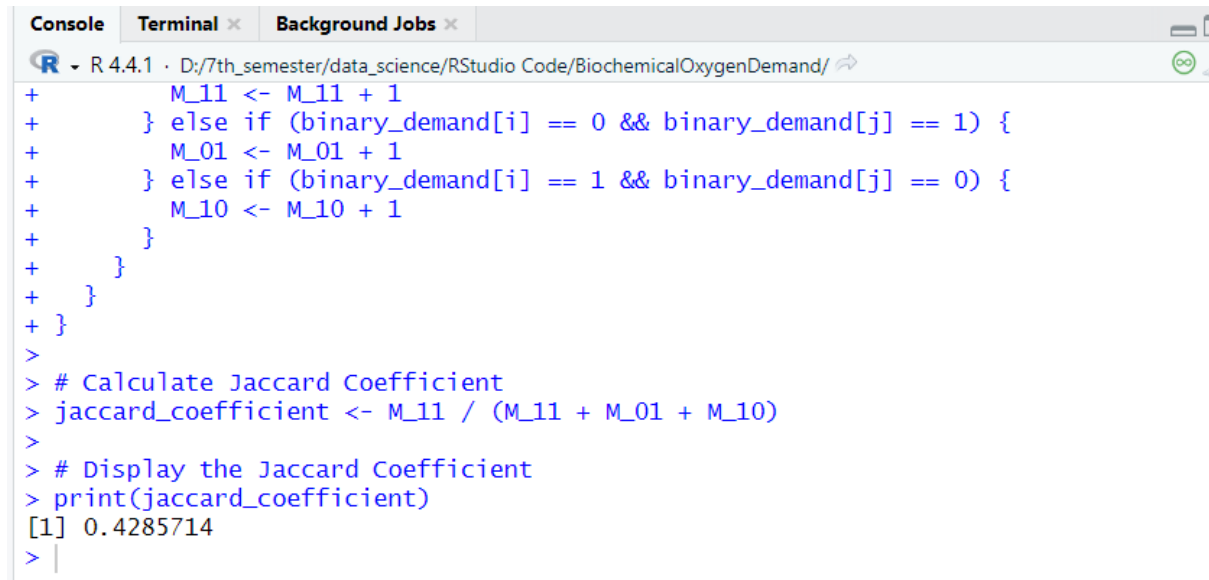

# Loop through each pair to count M11, M01, and M10
for (i in 1:n) {
  for (j in 1:n) {
    if (i != j) {
      if (binary_demand[i] == 1 && binary_demand[j] == 1) {
        M_11 <- M_11 + 1
      } else if (binary_demand[i] == 0 && binary_demand[j] == 1) {
        M_01 <- M_01 + 1
      } else if (binary_demand[i] == 1 && binary_demand[j] == 0) {
        M_10 <- M_10 + 1
      }
    }
  }
}
```

```
# Calculate Jaccard Coefficient

jaccard_coefficient <- M_11 / (M_11 + M_01 + M_10)

print(jaccard_coefficient)
```

Output



```
R 4.4.1 · D:/7th_semester/data_science/RStudio Code/BiochemicalOxygenDemand/
+     M_11 <- M_11 + 1
+   } else if (binary_demand[i] == 0 && binary_demand[j] == 1) {
+     M_01 <- M_01 + 1
+   } else if (binary_demand[i] == 1 && binary_demand[j] == 0) {
+     M_10 <- M_10 + 1
+   }
+ }
+ }
+ }
>
> # Calculate Jaccard Coefficient
> jaccard_coefficient <- M_11 / (M_11 + M_01 + M_10)
>
> # Display the Jaccard Coefficient
> print(jaccard_coefficient)
[1] 0.4285714
> |
```

Cosine Similarity

The cosine measure computes the *angle* between the two documents, which is **insensitive** to the absolute length of the document

Implementation

```
# Load the BOD dataset and save it to a CSV
write.csv(BOD, "BOD_Dataset.csv", row.names = FALSE)

#or load it instead
#load("BOD")

# Read the dataset back into R
BOD_data <- read.csv("BOD_Dataset.csv")

# Define a function to calculate cosine similarity
cosine_similarity <- function(x, y) {
```



```

dot_product <- sum(x * y)
magnitude_x <- sqrt(sum(x^2))
magnitude_y <- sqrt(sum(y^2))

# Calculate cosine similarity
if (magnitude_x == 0 || magnitude_y == 0) {
  return(NA) # Avoid division by zero
} else {
  return(dot_product / (magnitude_x * magnitude_y))
}
}

# Get the number of demands
n <- length(BOD_data$demand)

# Initialize a matrix to hold the cosine similarities
cosine_similarities <- matrix(NA, n, n)

# Calculate the cosine similarity for each pair
for (i in 1:n) {
  for (j in 1:n) {
    cosine_similarities[i, j] <-
cosine_similarity(BOD_data$demand[i], BOD_data$demand[j])
  }
}

# Print the cosine similarity matrix
cat("Cosine Similarity Matrix:\n")
print(cosine_similarities)

```

Output

```
> # Print the cosine similarity matrix
> cat("Cosine Similarity Matrix:\n")
Cosine Similarity Matrix:
> print(cosine_similarities)
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]    1    1    1    1    1    1
[2,]    1    1    1    1    1    1
[3,]    1    1    1    1    1    1
[4,]    1    1    1    1    1    1
[5,]    1    1    1    1    1    1
[6,]    1    1    1    1    1    1
```

Correlation

- Correlation measures the **linear relationship** between the attributes of the objects
 $y_k = ax_k + b$ for binary or continuous variable
- Measure of linear dependency b/w two variables x and y

Implementation

```
# Load the BOD dataset
```

```
BOD_data <- read.csv("BOD_Dataset.csv")
```

```
# Extract the Time and Demand columns
```

```
time <- BOD_data$Time
```

```
demand <- BOD_data$demand
```

```
# Calculate means
```

```
mean_time <- mean(time)
```

```
mean_demand <- mean(demand)
```

```
# Calculate covariance
```

```
covariance <- sum((time - mean_time) * (demand - mean_demand)) /  
length(time)
```

```
# Calculate standard deviations
```

```
std_dev_time <- sqrt(sum((time - mean_time)^2) / length(time))
```

```

std_dev_demand <- sqrt(sum((demand - mean_demand)^2) /
length(demand))

# Calculate Pearson correlation coefficient
correlation_coefficient <- covariance / (std_dev_time *
std_dev_demand)

# Print results
cat("Mean Time:", mean_time, "\n")
cat("Mean Demand:", mean_demand, "\n")
cat("Covariance:", covariance, "\n")
cat("Standard Deviation of Time:", std_dev_time, "\n")
cat("Standard Deviation of Demand:", std_dev_demand, "\n")
cat("Pearson Correlation Coefficient:", correlation_coefficient,
"\n")

```

Output

```

>
> # Print results
> cat("Mean Time:", mean_time, "\n")
Mean Time: 3.666667
> cat("Mean Demand:", mean_demand, "\n")
Mean Demand: 14.83333
> cat("Covariance:", covariance, "\n")
Covariance: 6.694444
> cat("Standard Deviation of Time:", std_dev_time, "\n")
Standard Deviation of Time: 1.972027
> cat("Standard Deviation of Demand:", std_dev_demand, "\n")
Standard Deviation of Demand: 4.227161
> cat("Pearson Correlation Coefficient:", correlation_coefficient, "\n")
Pearson Correlation Coefficient: 0.8030693
~ |

```

Clustering using the distance matrix

Load the BOD dataset and save it to a CSV

```
write.csv(BOD, "BOD_Dataset.csv", row.names = FALSE)
```

Read the dataset back into R

```
BOD_data <- read.csv("BOD_Dataset.csv")
```

```

# Binarize the demand data based on a threshold
threshold <- 15
BOD_data$binarized_demand <- ifelse(BOD_data$demand > threshold,
1, 0)

# Define the distance functions
euclidean_dist <- function(x, y) {
  sqrt(sum((x - y) ^ 2))
}

minkowski_dist <- function(x, y, r) {
  sum(abs(x - y) ^ r)^(1/r)
}

cosine_similarity <- function(x, y) {
  dot_product <- sum(x * y)
  magnitude_x <- sqrt(sum(x^2))
  magnitude_y <- sqrt(sum(y^2))

  if (magnitude_x == 0 || magnitude_y == 0) {
    return(NA) # Avoid division by zero
  } else {
    return(dot_product / (magnitude_x * magnitude_y))
  }
}

# Get the number of demands
n <- nrow(BOD_data)
demand_matrix <- as.matrix(BOD_data$demand)

# Calculate distance matrices

```

```

euclidean_matrix <- matrix(NA, n, n)
minkowski_matrix <- matrix(NA, n, n)
cosine_matrix <- matrix(NA, n, n)

# Calculate the distance matrices
for (i in 1:n) {
  for (j in 1:n) {
    # Using demand values for distance measures
    euclidean_matrix[i, j] <- euclidean_dist(demand_matrix[i, ,
drop = FALSE], demand_matrix[j, , drop = FALSE])
    minkowski_matrix[i, j] <- minkowski_dist(demand_matrix[i, ,
drop = FALSE], demand_matrix[j, , drop = FALSE], r = 3) # r can
be 1, 2, or 3
    cosine_matrix[i, j] <- cosine_similarity(demand_matrix[i, ,
drop = FALSE], demand_matrix[j, , drop = FALSE])
  }
}

# Print the distance matrices
cat("Euclidean Distance Matrix:\n")
print(euclidean_matrix)
cat("\nMinkowski Distance Matrix:\n")
print(minkowski_matrix)
cat("\nC cosine Similarity Matrix:\n")
print(cosine_matrix)

# Clustering using K-means based on Euclidean distance
set.seed(123) # For reproducibility
k <- 2 # Number of clusters

# Since K-means uses Euclidean distance, we will use the
original demand matrix
kmeans_result <- kmeans(demand_matrix, centers = k)

```

```
# Add cluster assignment to the original dataset
BOD_data$cluster <- kmeans_result$cluster
```

```
# Print the clustering result
cat("\nK-means Clustering Results:\n")
print(BOD_data)
```

Output

```
K-means Clustering Results:
> print(BOD_data)
  Time demand binarized_demand cluster
1    1    8.3                0        1
2    2   10.3                0        1
3    3   19.0                1        2
4    4   16.0                1        2
5    5   15.6                1        2
6    7   19.8                1        2
> |
```

References

- [1]. https://en.wikipedia.org/wiki/Euclidean_distance
- [2]. [https://en.wikipedia.org/wiki/Simple_matching_coefficient#:~:text=The%20simple%20matching%20coefficient%20\(SMC,and%20diversity%20of%20sample%20sets.&text=valu,e%201%2C%20and-,is%20the%20total%20number%20of%20attributes%20where%20A%20has,and%20B%20has%20value%200.](https://en.wikipedia.org/wiki/Simple_matching_coefficient#:~:text=The%20simple%20matching%20coefficient%20(SMC,and%20diversity%20of%20sample%20sets.&text=valu,e%201%2C%20and-,is%20the%20total%20number%20of%20attributes%20where%20A%20has,and%20B%20has%20value%200.)
- [3]. https://en.wikipedia.org/wiki/Jaccard_index